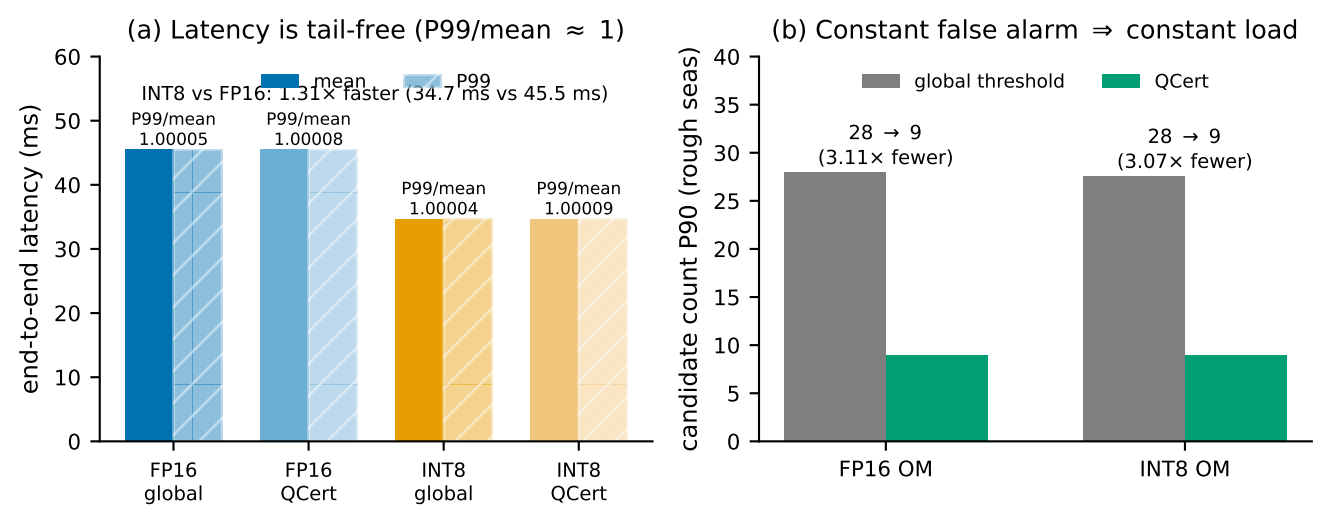
(a) Latency is tail-free (P99/mean ≈ 1)
mean
P99
INT8 vs FP16: 1.31× faster (34.7 ms vs 45.5 ms)
end-to-end latency (ms)
P99/mean 1.00005
P99/mean 1.00008
P99/mean 1.00004
P99/mean 1.00009
FP16 global
FP16 QCert
INT8 global
INT8 QCert
(b) Constant false alarm ⇒ constant load
global threshold
QCert
candidate count P90 (rough seas)
28 → 9 (3.11× fewer)
28 → 9 (3.07× fewer)
FP16 OM
INT8 OM